

REI204M: High Performance Computing
Industry Course Project
Orb - Forest Assessment

Emil Atli Þórhallsson
Samuel Caspar Gerbers

11. May 2025

GitHub Repository:

All code, data processing, and analysis scripts used in this report are available at:
<https://github.com/EmilAtli/Orb-project>

Introduction

With this project the goal is to see how we can translate the efforts from Meta in their project to predict tree heights based on high resolution satellite data [Tolan et al., 2024] to new data. In the paper the focus was on Los Angeles, USA and Sao Paulo, Brazil and we want to see if the model can also be used for Iceland's terrain. For that purpose we have tree data from Iceland, that we use for testing. A success of the project would mean, that the process of measuring trees in Iceland to watch the tree population can be made immensely easier and less time intensive.

Running the Model

Since the project is using a model using CUDA, making use of a high performance system will be a huge help. In the beginning we didn't have access to a system so we used Google Colab with its NVIDIA Tesla T4 runtime, you can see the procedure in the ORB.ipynb Jupyter notebook in the zip or on our Github. Later we switched to using the Elja system to run the project.

Preparing the environment for Elja

We start by accessing the system using our ssh key:

```
1 | ssh -i ~/.ssh/elja scg6@elja.hi.is
```

Then we set up the folder structure:

```
1 | mkdir ORB
2 | cd ORB
```

Now we are on the login node and that is where we prepare the virtual environment for Python to be able to run the script.

First to make sure everything is consistent we also want to use the specific Python version that we want to use for the project which is Python 3.11 and with `module avail` in the system we can see that there is even a module with Python 3.11.3 which we can use:

When we wanted to run the program we had a problem with pytorch needing to have a compiler, which it didn't have but we could fix it by adding a dummy compiler to the inference.py script inside the git repository. To add it you can run the following python script or add the block in the python script manually:

```

1 import fileinput
2 # Inserts a dummy compiler, since torch v2.0.1 doesn't have a
  compiler.
3
4 # block to insert.
5 block = '''if not hasattr(torch, "compiler"):
6     class DummyCompiler:
7         @staticmethod
8         def disable(*args, **kwargs):
9             def decorator(func):
10                 return func
11             return decorator
12     torch.compiler = DummyCompiler()
13 '''
14
15 # Open inference.py for in-place editing.
16 with fileinput.input('./inference.py', inplace=True) as file:
17     for line_number, line in enumerate(file, 1):
18         print(line, end='')
19         # After line 8 (which is "import torch"), block is inserted
20
21         if line_number == 8:
22             print(block, end='')

```

So to execute the script you can simply create a file using vim and then copy the data there and execute the file:

```

1 vim fix_inference.py
2 # paste the code above inside, save it and run the script:
3 python3 fix_inference.py

```

Now we have to prepare the data we want to use for the inference so the tree data that we were provided. I still had it on my local machine so we can use scp to copy it from my local machine to the remote HPC system, so we execute the following command from our local machine:

```

1 scp -i ~/.ssh/elja trees.zip scg6@elja.hi.is:/jotunn/home/scg6/ORB/
  HighResCanopyHeight

```

After that we also have to prepare the csv file that correctly uses the new data set, for that purpose we made a short python script that we executed on our local machine after unzipping the tree data there. So we:

```

1 unzip trees.zip

```

And then we can use the following make_csv.py script to create the new csv:

```

1 import os
2 import csv
3
4 file_list = sorted(os.listdir('./output_images/'))
5
6 with open("data.csv", "w", newline="") as csvfile:

```

```

7 writer = csv.writer(csvfile)
8 writer.writerow(["", "index", "image"]) # header with leading
  comma
9 for idx, filename in enumerate(file_list):
10     index = int(filename[:-4])
11     writer.writerow([idx, index, filename])

```

Now we also want to send this csv file to Elja:

```

1 scp -i ~/.ssh/elja data.csv scg6@elja.hi.is:/jotunn/home/scg6/ORB/
  HighResCanopyHeight

```

Finally before executing the the programm we also need to unzip the trees.zip file on Elja:

```

1 unzip trees.zip

```

Running the model on Elja

Now everything is prepared to actually execute the model on Elja, to make sure we use the right partition and we understand how it is structured to write the batch script to submit the job we used `sinfo` and `scontrol show node`:

```

[scg6@elja-irhpc ~]$ sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
Jotunn      up 1-00:00:00      4  idle compute-[1-4]
Jotunn-GPU  up 2-00:00:00      1  alloc gpu-2

```

Figure 3: Result of sinfo

```

[scg6@elja-irhpc ~]$ scontrol show node gpu-2
NodeName=gpu-2 Arch=x86_64 CoresPerSocket=32
CPUAlloc=128 CPUEffctv=128 CPUTot=128 CPUload=1.01
AvailableFeatures=(null)
ActiveFeatures=(null)
Gres=(null)
NodeAddr=gpu-2 NodeHostName=gpu-2 Version=23.11.10
OS=Linux 4.18.0-553.51.1.el8_10.x86_64 #1 SMP Wed Apr 30 20:24:04 UTC 2025
RealMemory=252000 AllocMem=252000 FreeMem=243734 Sockets=2 Boards=1
State=ALLOCATED ThreadsPerCore=2 TmpDisk=0 Weight=1 Owner=N/A MCS_label=N/A
Partitions=Jotunn-GPU
BootTime=2025-05-09T12:36:22 SlurmdStartTime=2025-05-09T12:38:16
LastBusyTime=2025-05-11T11:58:18 ResumeAfterTime=None
CfgTRES=cpu=128,mem=252000M,billing=128
AllocTRES=cpu=128,mem=252000M
CapWatts=n/a
CurrentWatts=0 AveWatts=0
ExtSensorsJoules=n/a ExtSensorsWatts=0 ExtSensorsTemp=n/a
[scg6@elja-irhpc ~]$

```

Figure 4: Result of scontrol show node

Using that info we can write our batch script `submit.sh` to submit the job, here we can then simply activate the python environment we created before:

```

1 #!/bin/bash
2 #SBATCH -J orb
3 #SBATCH --partition=Jotunn-GPU

```

```

4 #SBATCH --nodes=1
5 #SBATCH --ntasks-per-node=1
6 #SBATCH --cpus-per-task=4
7 #SBATCH --mem=32G
8 #SBATCH --time=01:00:00
9 #SBATCH --mail-user=scg6@hi.is
10 #SBATCH --mail-type=end
11
12 module load Python/3.11.3
13 source ../orb_env/bin/activate
14
15 python inference.py --csv ./data.csv --image_dir ./output_images/
    --name output_inference

```

Then we can actually submit it and see it in the queue:

```

(orb_env) [scg6@elja-lrhpc HighResCanopyHeight]$ sbatch submit.sh
Submitted batch job 1351611
(orb_env) [scg6@elja-lrhpc HighResCanopyHeight]$ squeue

```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NOELIST(REASON)
1349798	48cpu_192	DecJan20	haa53	R	3-19:28:10	1	compute-17
1351345	48cpu_192	md_NPC	aptorres	PD	0:00	1	(Priority)
1350502	64cpu_256	EAG4	mynent	R	2-13:26:17	1	compute-63
1351611	Jotunn-GP	orb	scg6	R	0:01	1	gpu-2

Figure 5: Submitting the job

Since we specified we get a notification at the end of the job via mail, we get the runtime from there:

```
Slurm Job_id=1351611 Name=orb Ended, Run time 00:07:13, COMPLETED, ExitCode 0
```

Figure 6: finished job

It only took 07:13 to complete the job, showing the how useful it is to use a HPC system like Elja.

Since a lot of of the trees are in the same areas, they are combined in one picture instead of having individual pictures for each tree. Now we have all the resulting trees in the output_inference folder, which we can zip and then send back to our local machine to do the analysis of the data there:

```

1 (on elja) zip -r output_inference.zip output_inference/
2 (on local machine) scp -i ~/.ssh/elja scg6@elja.hi.is:/jotunn/home/
    scg6/ORB/HighResCanopyHeight/output_inference.zip /home/samu/
    Documents/Iceland/Uni/Orb/submit/

```


Combining the data

The zip trees.zip contains the output.gpkg file which has all the information about the trees including the position of the tree in the image and also the ground truth height of the tree. Using this file we can easily combine the ground truth with the predictions of the model.

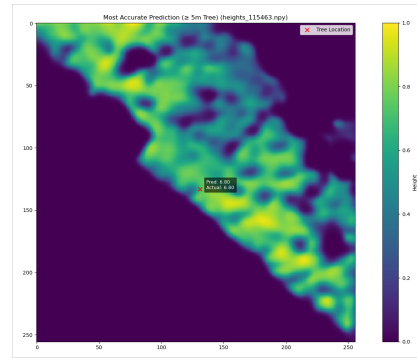
To combine it we can simply add a column describing the predicted height, this can be done with a simple python script by checking, the location of the height map that was generated by the model for the height that we predicted for each tree and then adding that number to the row of the tree:

```
1 import geopandas as gpd
2 import numpy as np
3 import os
4 from scipy.ndimage import map_coordinates
5
6 input_gpkg = "./output.gpkg"
7 input_layer = "output"
8 npy_folder = "./inference/"
9 output_gpkg = "./output_with_predictions.gpkg"
10 output_layer = "predicted_output"
11
12 gdf = gpd.read_file(input_gpkg, layer=input_layer)
13
14 def get_precise_height(image_name, px, py):
15     npy_name = "heights_" + image_name.replace(".png", ".npy")
16     npy_path = os.path.join(npy_folder, npy_name)
17     if not os.path.exists(npy_path):
18         return None
19     arr = np.load(npy_path)
20     if 0 <= px < arr.shape[1] and 0 <= py < arr.shape[0]:
21         coords = np.array([[py], [px]])
22         return float(map_coordinates(arr, coords, order=1, mode='
nearest')[0])
23     return None
24
25 gdf["predicted_height"] = gdf.apply(
26     lambda row: get_precise_height(row["image"], row["px"], row["py"]
27     ),
28     axis=1
29 )
30 gdf.to_file(output_gpkg, layer=output_layer, driver="GPKG")
31 print("done")
```

And we can also look at some specific examples of our prediction here is one example where it performed good and one where it performed bad:

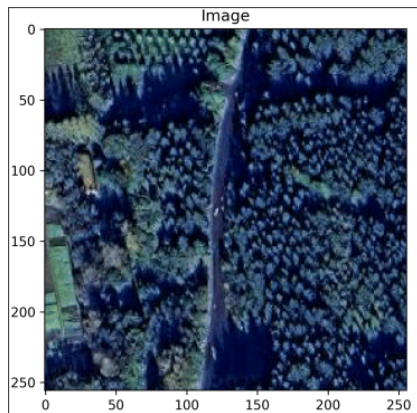


(a) Satellite Picture

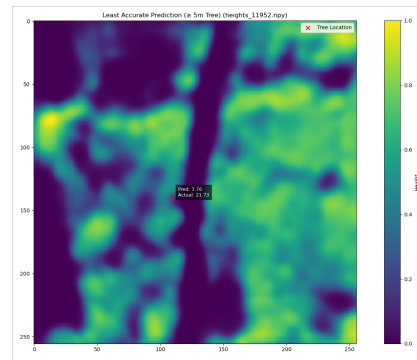


(b) Prediction

Figure 7: Good performing sample



(a) Satellite Picture



(b) Prediction

Figure 8: Bad performing sample

Assessing the Performance

Linear Regression




Figure 9: Result of Linear Regression

The figure above shows the result of applying a linear regression to compare the predicted tree heights from the model with the ground truth measurements from the entire dataset. Each dot represents an individual tree, with the predicted height on the x-axis and the actual measured height on the y-axis.

The black line is the fitted regression line, showing the general trend between predictions and reality. The blue shaded area around the line represents the 95% confidence interval, which tells us how certain the model is about the average trend. The wider orange band shows the 95% prediction interval, which indicates the expected range for individual predictions.

We calculated an R^2 score of 0.40, meaning the model explains about 40% of the variance in the data. The root mean squared error (RMSE) is 1.97 meters, which means that, on average, the model's predictions are off by nearly 2 meters. Based on this result, the model shows a general trend but struggles to make precise predictions for individual trees when using the full dataset.

Table 1: Regression results for each species, sorted by R^2 score.

Species	R^2	RMSE
Hvítgreni	0.92	0.58
Gráö lur	0.86	1.27
Reynir	0.83	0.38

Continued on next page

Continued from previous page

Species	R^2	RMSE
Broddgreni	0.70	1.57
Lindifura	0.60	0.57
Hengibirki	0.53	1.66
Rússa- og síberíulerki	0.51	1.90
Sitkagreni	0.51	1.70
Selja	0.50	1.22
Stafafura	0.49	1.78
Sitkagreni og -bastarður	0.47	2.07
Rússalerki	0.45	1.65
Alaskaösp	0.42	3.06
Rauðgreni	0.41	2.41
Gulvíðir	0.38	1.22
Reyniviður	0.37	2.02
Sitkaölur	0.30	0.75
Bergfura	0.29	1.55
Ilmbjörk	0.26	1.51
Fjallapöll	0.25	0.24
Evrópulerki	0.25	0.59
Blágreni	0.17	1.19
Skógarfura	0.15	1.91
Víðir	0.13	1.48
Viðja	0.10	1.75
Alaskavíðir	0.10	1.47
Ilmbjörk - sjálfsáð	0.09	1.35
Jörfavíðir	0.07	0.45
Fjallapínur	0.00	0.91

Table 1 shows how the model performed on each tree species, based on R^2 and RMSE. Looking at the results, it's clear that the model works best for *Hvítgreni*, *Gráölur*, and *Reynir*, which all have R^2 values over 0.8. This means the model does a pretty good job of predicting tree heights for these species. The RMSE for these species is also quite low, which adds to our confidence in the results.

Further down the list, the R^2 values drop off quite a bit. For many of the other species, the model struggles to make reliable predictions. Some of the R^2 scores are even close to zero, which basically means the model isn't doing much better than random guessing for those species.

It's also worth mentioning that a few species were skipped in the analysis because there was only one sample available. These species were *Silfurreynir*, *Álmur*, *Douglasgreni*, *Kjarrölur*, *Fura*, *Heggur*, and *Fjallafura*. With just one data point, it isn't possible to fit a meaningful regression model, so these were left out.

All in all, it seems like the model shows potential for certain species, but there's still a lot of room for improvement if we want it to work well across the board.

Improved Sampling with Neighborhood Maximum

In the previous analysis, the predicted height for each tree was sampled directly at the pixel location provided in the ground truth data. However, this assumes that the provided pixel coordinates perfectly align with the actual treetop in the prediction map, which is often not the case. To account for this, we updated the method to search for the highest predicted value in a local neighborhood around each tree location.

After testing different window sizes, ranging from 3x3 to 21x21 pixels, we found that using a 21x21 window (covering approximately 10.5 meters by 10.5 meters) produced the best results. This window size likely captures the actual treetop even if the provided coordinates are slightly off. The improved regression analysis using this method achieved an R^2 score of 0.48 and an RMSE of 1.84 meters, slightly better than the previous approach.

Figure 10 shows the updated regression result with the improved sampling method. The results suggest that allowing the model to search for the local maximum in the area around each tree provides a more reliable comparison to the ground truth.




Figure 10: Linear regression result using the maximum value in a 21x21 pixel window around each tree location.

Table 2: Regression results per species after applying neighborhood maximum sampling (21x21 window).

Species	R^2	RMSE
Lindifura	0.92	0.25
Hvítgreni	0.89	0.65
Reynir	0.83	0.38
Rússalerki	0.66	1.30
Rússa- og síberíulerki	0.62	1.67
Alaskaösp	0.57	2.63
Stafafura	0.55	1.68
Gulvíðir	0.54	1.04
Rauðgreni	0.54	2.14
Sitkagreni og -bastarður	0.51	2.00
Gráölur	0.50	2.37
Víðir	0.49	1.13
Sitkagreni	0.49	1.74
Sitkaölur	0.47	0.66
Reyniviður	0.41	1.94
Ilmbjörk	0.36	1.39
Broddgreni	0.35	2.33
Bergfura	0.34	1.51
Hengibirki	0.29	2.03
Evrópulerki	0.24	0.59
Viðja	0.19	1.66
Skógarfura	0.15	1.91
Ilmbjörk - sjálfsáð	0.11	1.33
Jörfavíðir	0.09	0.45
Fjallaþöll	0.07	0.27
Alaskavíðir	0.06	1.50
Blágreni	0.05	1.27
Fjallaþinur	0.02	0.91
Selja	0.01	1.71

Table 2 shows the updated regression results after applying the neighborhood maximum sampling method using a 21x21 pixel window. Compared to the previous analysis, which relied on single-pixel sampling, this approach generally improved the R^2 scores for most species. Notably, *Lindifura* now achieves the highest R^2 of 0.92 with a low RMSE of 0.25 meters. This is a clear improvement over the earlier results, where *Hvítgreni* had the best performance.

Another interesting thing is that several species, such as *Rússalerki*, *Rússa- og síberíulerki*, and *Alaskaösp*, show higher R^2 values than before, suggesting that the neighborhood search helps the model better capture the true tree heights for these species.

However, the improvements are not universal. Some species, especially those with lower R^2 values in the earlier table, remain challenging for the model. While the neighborhood maximum method provides a clear overall benefit, it does not fully resolve the limitations for all species.

Overall, looking at both tables side by side, it seems pretty clear that letting the model search for the highest value in the area around each tree helps. This is especially true for species where the exact tree location might not line up perfectly with the prediction map.

Distribution of Predictions and Errors

To get a better understanding of how our model performs across different parts of Iceland, we visualized the spatial distribution of both the predicted tree heights and the ground truth heights. This also allowed us to generate an error map, showing how far off the predictions are compared to the actual measured tree heights.

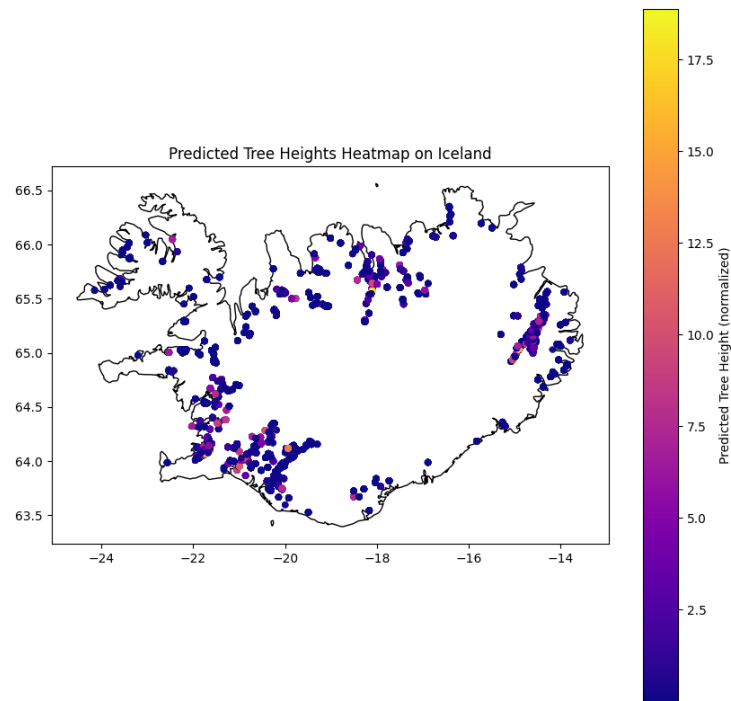


Figure 11: Predicted tree heights plotted on Iceland's outline.

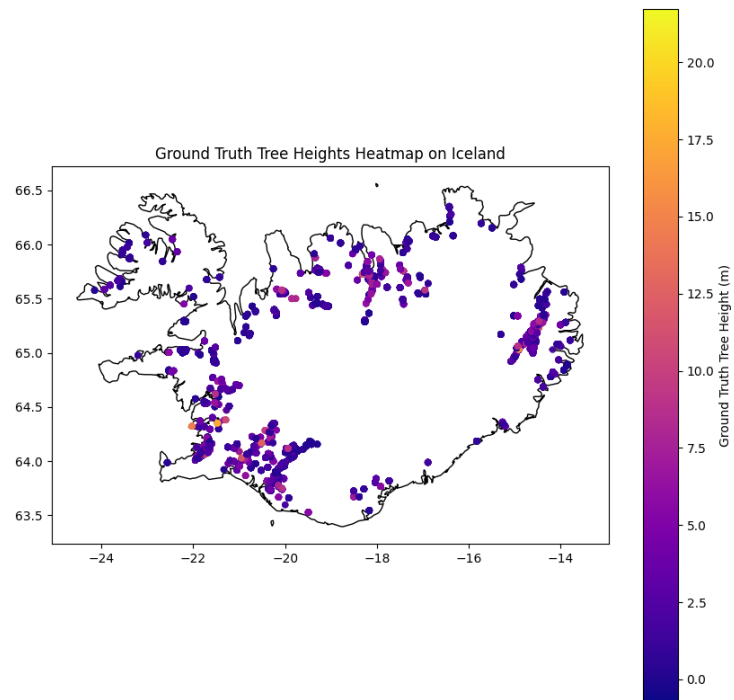


Figure 12: Ground truth tree heights plotted on Iceland's outline.

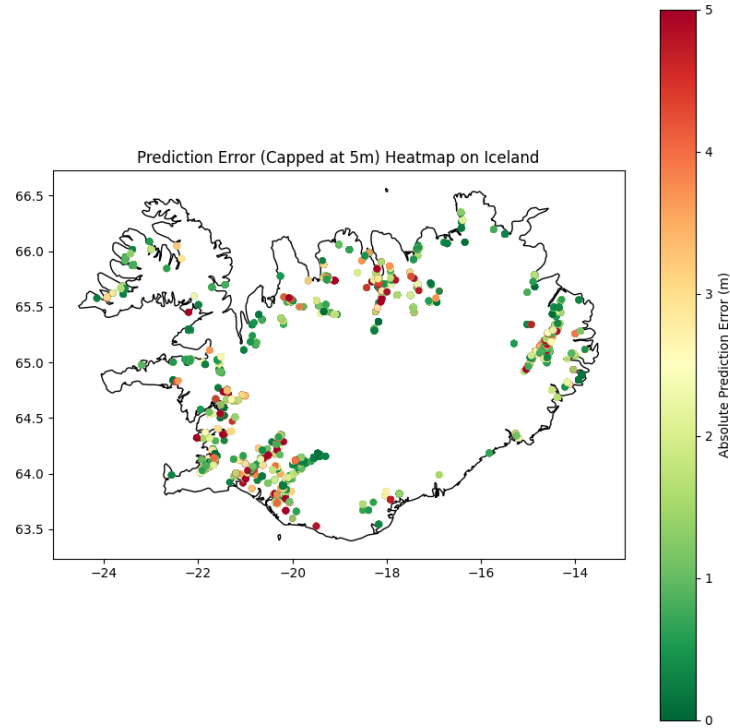


Figure 13: Prediction error heatmap, showing how far off the predictions are compared to ground truth. Green areas indicate good alignment, while red areas show larger errors. The error scale is capped at 5 meters to highlight the typical error range.

Looking at these maps, it becomes clear that while the general distribution of predicted and actual tree heights matches reasonably well, there are localized areas with higher errors. The error map shows that although the model performs well in many areas (green dots), there are clusters where the predictions are off by several meters (yellow to red dots). This spatial overview helps us identify regions where the model might need improvement or where the input data might be less reliable.

References

[Tolan et al., 2024] Tolan, J., Yang, H.-I., Nosarzewski, B., Couairon, G., Vo, H. V., Brandt, J., Spore, J., Majumdar, S., Haziza, D., Vamaraju, J., Moutakanni, T., Bojanowski, P., Johns, T., White, B., Tiecke, T., and Couprie, C. (2024). Very high resolution canopy height maps from rgb imagery using self-supervised vision transformer and convolutional decoder trained on aerial lidar. *Remote Sensing of Environment*, 300:113888.